


A woman with dark hair tied back, wearing a light-colored knit sweater, is looking down at a black smartphone she is holding with both hands. The background is a soft-focus indoor setting with a white cup of coffee on a saucer, a potted snake plant, and an open notebook on a wooden table. The lighting is warm and natural, suggesting a morning or afternoon setting.

Fable: the companion bet

How a mental wellness app keeps growing
when its users get better

Product Leadership Certification · Final Project

Product School · August 2026 cohort · Antje Barth

Fable's real competitor is the user deciding they are done

Playing to Win

Choice	Our answer
Winning aspiration	The first thing you talk to when something feels off, because it has known you for two years.
Where to play	Existing users, plus recent graduates, meaning people who got better and stopped using it, 25 to 45, US English first. Everyday wellbeing, not just acute episodes. Not clinical care, not general assistance.
How to win	Compounding personal context inside a clinically grounded companion. After a year, switching means starting over with a stranger.
Capabilities	Build the memory engine and clinical conversation design. Partner for the model, clinical advisors and crisis escalation. World-class at safety.
Management systems	Felt-understood score and 26-week retention. Weekly safety review with halt authority, monthly clinical audit, safety-eval ship gate.

The hard no

No attachment optimization, friendship or romance simulation, time-in-app targets, or dependency mechanics. Attachment destroys the clinical trust that is our only durable moat.

OKRs

KR1, share of users still here at 26 weeks: 12% to 30%. KR2, share of conversations people rate 4 or better on 'did this feel understood': 38% to 65%, clinically audited. KR3, users who got better and left, coming back for a second conversation: 5% to 18%.

Three Rocks, four cuts, and the cuts are where the strategy is

Now, committed

The three Rocks, meaning the bets that have to land or the strategy does not.

Memory-grounded check-ins: the thing competitors cannot copy. Drives KR2. *Memory & Personalization, Conversation Design.*

Crisis mode: what makes it safe to run this at all, and the safe way back for someone who left. Drives KR3. *Safety & Clinical, Platform.*

Sub-2-second load: if it takes four seconds to open at 3am, the promise is already broken. Drives KR1. *Platform, Engineering.*

Performance as a Rock, deliberately

Speed usually gets filed as polish, not strategy. Here it is a Rock because the product's whole promise is being the first thing you reach for, and nothing that takes four seconds is the first thing you reach for.

Later, bets not commitments

Voice, web and watch surfaces. Spanish and Portuguese markets. A supporter track for the people around a struggling user, plus clinician-ready export.

What we cut

- **Daily streaks:** the hard no in feature form
- **Content library expansion:** invests further in the very thing users finish with
- **Social sharing:** mental-health progress is not social currency
- **Fable for Teams:** employer reporting conflicts with the hard no and the privacy promise. Revisit after KR2.

Sequencing the Next queue

WSJF, which ranks work by value and urgency divided by effort, puts the pricing pilot top of the queue at 6.00, mostly because it removes risk. Price gate set at 30 November rather than waiting for the retention read.

A hard no only works if a ticket can be tested against it

What we own

We own the companion relationship end to end. A shared platform with product teams on top. Memory, conversation design and safety are built once, centrally; three PMs own outcomes that use them. We optimize for coherence and accept slower velocity in the shared layer, because a companion that behaves like two products is not a companion. The rule that settles the overlap: whoever owns the screen decides where and when a check-in appears, and whoever owns the check-in decides what it says.

How we decide

Metric owner decides. Capability owner can veto. Safety and the hard no stop anything outright. Not overridable. No cross-team conflict stays open more than three business days.

The seam this charter had to close

Retention's fastest levers and the hard no look identical in a ticket. Both are a notification. The line: a check-in is **earned return** if a user-specific signal triggers it, and **manufactured dependency** if elapsed time or a frequency target triggers it.

What success looks like

26-week retention 12% to 30%. Felt-understood 4+ share 38% to 65%. Graduate second-conversation rate 5% to 18%. Zero unhandled safety incidents. Miss retention and felt-understood together and we replan the check-in roadmap before next quarter.

Not a growth bet. A repair of unit economics sitting within cents of cost of capital

Where Fable is today

For every **\$1.00** we spend acquiring someone, we get **\$1.05** back, and it takes 25 months. That is breaking even, not earning. We need **\$9.57** a month per user just to break even, and we charge **\$9.99**, so we clear the line by **42 cents**. If people leave a little faster than assumed, breakeven rises to **\$10.20** and we lose **21 cents** on each one. Nobody has measured which world we are in.

What the bet does

A **\$1.0M** build takes that to **\$1.82** back on every **\$1.00**, paying back in 10 months, worth about **\$11.6M** a year. It drops the break-even point to **\$6.40** a month. The **\$3.59** of room above break-even is the real product, not the return.

The ask

\$1.0M, six engineers for two quarters, gated on the felt-understood score hitting **55%** by 31 December.

Stage	LTV:CAC
Today	1.05:1
After the bet	1.82:1
Sustainable <i>(needs roughly 59% retention, unreachable here)</i>	3:1

Seven kill criteria. Here is one.

If felt-understood 4+ share has not reached **55%** by **31 December 2026**, we do not commit the Q1 2027 headcount: six engineers, about **\$375K** per quarter, return to crisis mode. The gate that matters most checks, at month 12, how long people actually stay once they pass six months, which carries **74%** of the value.

A strategy is only as real as the number that would end it

Friction

The hard no took two minutes to write and four modules to make usable. Module 2 made me cut real revenue. Module 3 showed the rule could not be applied to a ticket. Module 4 caught me about to concede the loophole myself: time since last open is a user-specific signal.

Learning: good economics, wrong bet

Good unit economics don't make a bet worth funding. The case I reviewed had better numbers than mine. Still a no. The ceiling was **\$414K** against a 4.2M-user business.

Learning: the wrong assumption scared me

The assumption I feared wasn't the one carrying the case. Being 30% wrong about how strongly feeling understood drives retention still leaves the bet worth funding. But how long people stay once they pass six months, a number nobody has measured, carries **74%** of the value.

The aha

The strategy felt done when all the pieces lined up, but nothing in it could be proven wrong. Writing kill criteria changed that. Judgment only travels if someone who isn't me can act on it.