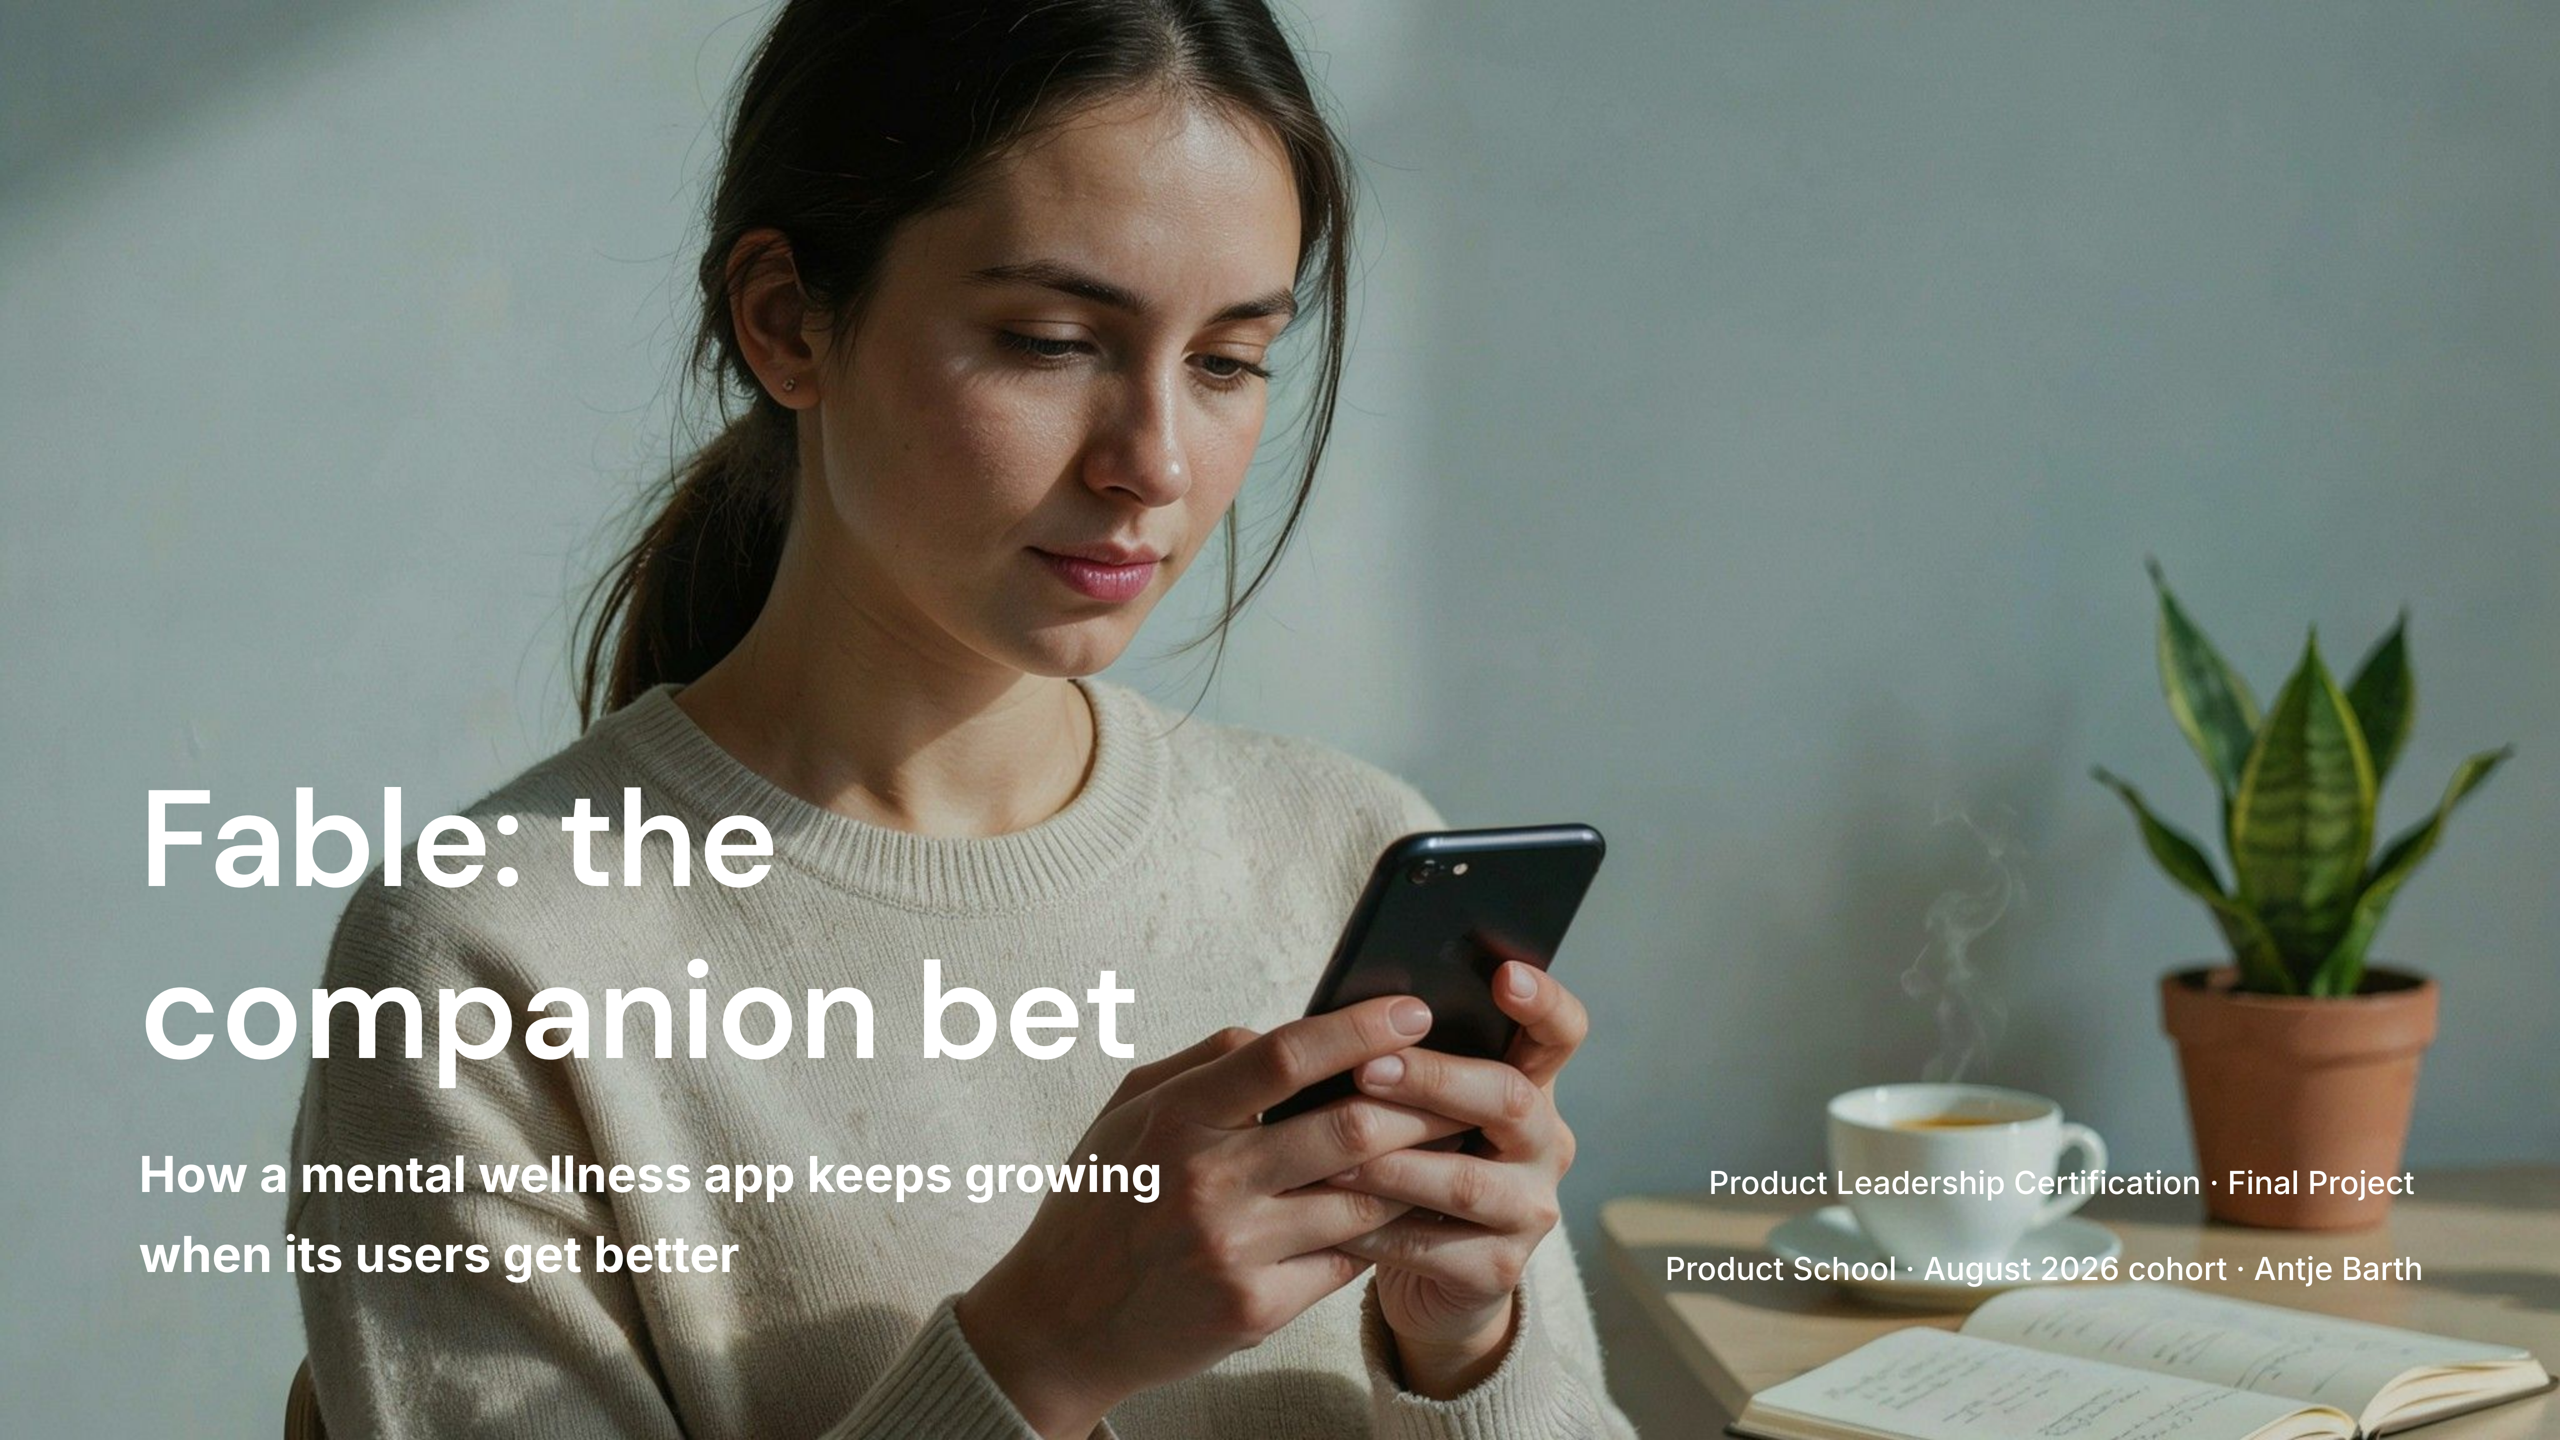

Fable: the companion bet

How a mental wellness app keeps growing when its users get better

Product Leadership Certification · Final Project

Product School · August 2026 cohort · Antje Barth

SLIDE 1 · PRODUCT STRATEGY

Fable's real competitor is the user deciding they are done

The diagnosis

Adults resolve anxiety in 90 to 120 days, then stop opening the app. 4.2M users, growth flat. Engagement mechanics argue with a user who is right.

The bet

A clinically grounded AI companion whose value compounds through longitudinal personal memory, winning the everyday "something feels off" moment.

How we win

ChatGPT starts cold. Calm and Headspace hold attention, not history. After a year with Fable, switching means starting over with a stranger.

The hard no

No attachment optimization, friendship or romance simulation, time-in-app targets, or dependency mechanics. Attachment destroys the clinical trust that is our only durable moat.

OKRs

26-week retention: 12% to 30%. Felt-understood 4+: 38% to 65% (clinically audited). Dormant graduates completing a second conversation: 5% to 18%.

Three Rocks, four cuts, and the cuts are where the strategy is

Now, committed

Memory-grounded check-ins: the differentiator. Drives KR2.

Crisis mode: the license to operate and the dormant graduate's safe return. Drives KR3.

Sub-2-second load: the promise made physical. Drives KR1.

Performance as a Rock, deliberately

For the first thing users open when something feels off, a 4-second cold start breaks the promise at the front door.

What we cut

- **Daily streaks:** the hard no in feature form
- **Content library expansion:** deepens the value prop users graduate out of
- **Social sharing:** mental-health progress is not social currency
- **Fable for Teams:** employer reporting conflicts with the hard no and the privacy promise. Revisit after KR2.

Sequencing the Next queue

WSJF puts the pricing pilot first at 6.00 on risk reduction. Price gate set at 30 November rather than waiting for the retention read.

SLIDE 3 · TEAM CHARTER

A hard no only works if a ticket can be tested against it

What we own

We own the companion relationship end to end. Platform plus Product: memory, conversation design and safety are centralized, and three PMs own outcomes on top. We optimize for coherence and accept slower velocity in the shared layer, because a companion that behaves like two products is not a companion. Crux rule where surfaces meet: surface owners own the frame, the check-in layer owns the content.

How we decide

Metric owner decides. Capability owner can veto. Safety and the hard no stop anything outright. Not overridable. No cross-team conflict stays open more than three business days.

The seam this charter had to close

Retention's fastest levers and the hard no look identical in a ticket. Both are a notification. The line: a check-in is **earned return** if a user-specific signal triggers it, and **manufactured dependency** if elapsed time or a frequency target triggers it.

What success looks like

26-week retention 12% to 30%. Felt-understood 4+ share 38% to 65%. Graduate second-conversation rate 5% to 18%. Zero unhandled safety incidents. Miss retention and felt-understood together and we replan the check-in roadmap before next quarter.

SLIDE 4 · FINANCIAL MODEL

Not a growth bet. A repair of unit economics sitting within cents of cost of capital

Where Fable is today

Breakeven ARPU is **$9.57** against the **$9.99** we charge. We clear it by **42 cents** at a 4% churn tail, and miss by **21 cents** at a 5.2% tail, where breakeven is **$10.20**.

What the bet does

A **$1.0M** build lowers breakeven ARPU to **$6.40**. That **$3.59** of headroom is the real product, not the return. Payback shortens from 25 to 10 discounted months, worth about **$11.6M** a year.

The ask

$1.0M, six engineers for two quarters, gated on **55%** by 31 December.

Stage	LTV:CAC
Today	1.05:1
After the bet	1.82:1
Sustainable *(needs roughly 59% retention, unreachable here)*	3:1

Seven kill criteria. Here is one.

If felt-understood 4+ share has not reached **55%** by **31 December 2026**, we do not commit the Q1 2027 headcount: six engineers, about **$375K** per quarter, return to crisis mode. The gate that matters most reads the churn tail at month 12, which carries **74%** of the value.

SLIDE 5 · INDIVIDUAL INSIGHTS

A strategy is only as real as the number that would end it

Friction

The hard no took two minutes to write and four modules to make usable. Module 2 made me cut real revenue. Module 3 showed the rule could not be applied to a ticket. Module 4 caught me about to concede the loophole myself: time since last open is a user-specific signal.

Learning: good economics, wrong bet

Good unit economics don't make a bet worth funding. The case I reviewed had better numbers than mine. Still a no. The ceiling was **$414K** against a 4.2M-user business.

Learning: the wrong assumption scared me

The assumption I feared wasn't the one carrying the case. A **30%** miss on the transfer coefficient still approves. A 4% churn tail nobody has measured carries **74%** of the value.

The aha

The strategy felt done when the cascade cohered, but nothing in it could be proven wrong. Writing kill criteria changed that. Judgment only travels if someone who isn't me can act on it.